


Lecture #15

- Priority Queues
 - Heaps
- HeapSort

Priority Queues



Priority Queue

Why should you care?

Priority Queues are used to prioritize different types of data for processing based on their importance.

They're used in:

Call answering centers

Turn-by-turn navigation software

Operating Systems (thread mgmt.)

Network packet routing

So pay attention!

Why
should
I care?




Priority Queues

A **priority queue** is a special type of queue that allows us to keep a prioritized list of items.

In a **priority queue**, each item you **insert** into the queue has a "**priority rating**" indicating how important it is.

Any time you **dequeue** an item from a priority queue, it always dequeues the item with the **highest priority** (instead of just the first item inserted).

Example: If I have a queue of patients in the emergency room, I don't just take the next patient in line, I take the one who has the most severe injuries.



Priority Queues

A **Priority Queue** supports three operations:

- Insert a new item into the queue
- Get the value of the highest priority item
- Remove the highest priority item from the queue

When you define a Priority Queue, you must specify how to determine the priority of each item in the queue, e.g.

Priority = amount of blood lost + number of cuts


You must then design your PQ data structure/algorithms so you can efficiently retrieve the highest-priority item.

Priority Queues

Question: What data structures can we use to implement a priority queue?

Let's make it easier... What if we have just a limited set of priorities, e.g.: **high**, **medium** **low**?

Hint: Think of an airport ticket line with **first**, **business** and **coach** (cattle) class...



Right - we can use $n=3$ linked lists, one for each priority level.

To obtain the highest-priority item, always take the first item from the highest priority, non-empty list.

Priority Queues

Question: Ok, but what data structure should we use if we have a huge number of priorities? Hmmm!

The **HEAP** data structure is one of the most efficient ones we can use to implement a Priority Queue.



The heap data structure uses a special type of **binary tree** to hold its data.

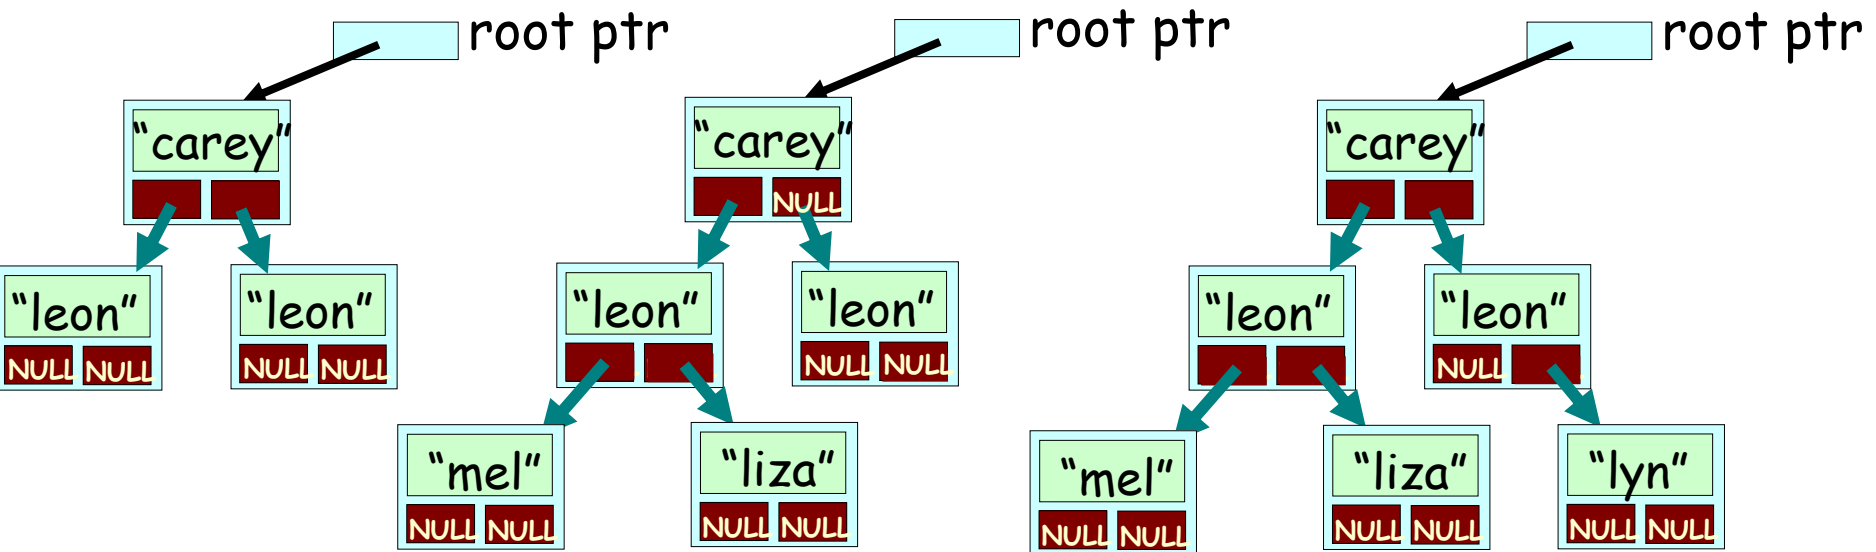
As we'll see, while a heap does use a binary tree to store its data, a heap is **NOT** a **binary search tree**.

All Heaps Use a "Complete" Binary Tree

A **complete binary tree** is one in which:

- The top $N-1$ levels of the tree are completely filled with nodes
- All nodes on the bottom-most level must be as far left as possible (with no empty slots between nodes!)

Is it complete (note: these examples are slightly different from the PPT slides)?



Answers: Left - yes, Middle - yes, Right - no - Leon's right child is not filled, yet his left child is present, so our bottom row is not completely filled as far left as possible.

Heaps... of...

A heap is a special type of **complete binary tree** (it's **not** a binary search tree).

There are two types of heaps, **minheaps** and **maxheaps**:

Maxheap:

1. Quickly insert a new item into the heap
2. Quickly retrieve the **largest** item from the heap

Minheap:

1. Quickly insert a new item into the heap
2. Quickly retrieve the **smallest** item from the heap

The Maxheap

A **maxheap** is a binary tree which follows these rules:

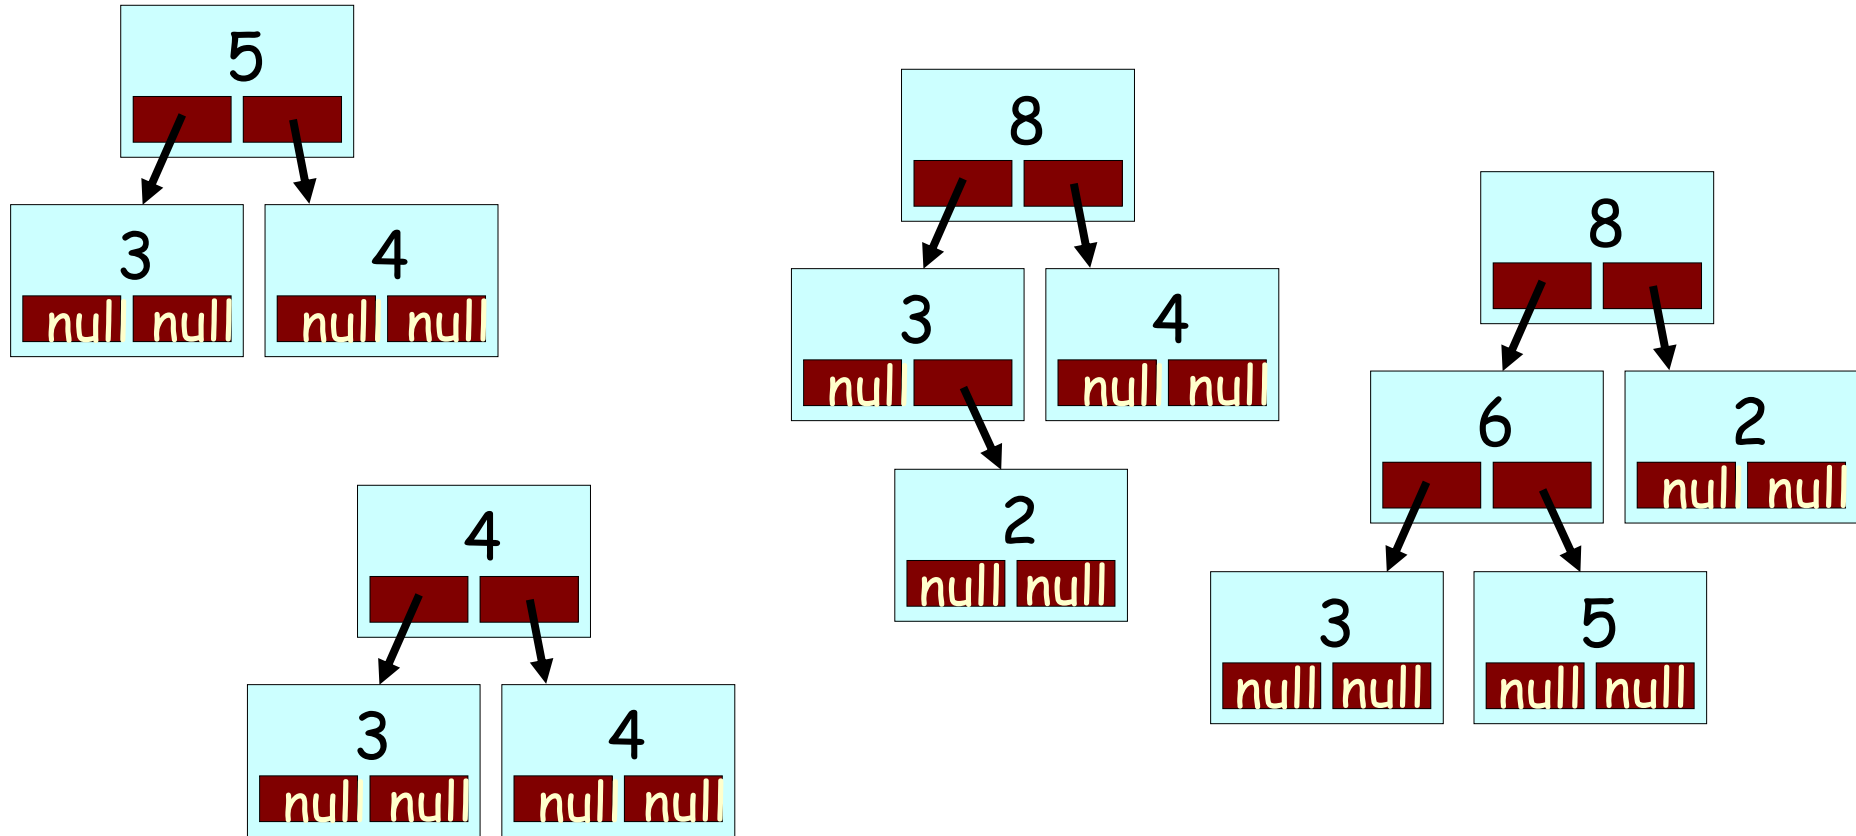
1. The value contained by a node is **ALWAYS GREATER THAN OR EQUAL TO** the values of the node's children.
2. The tree is a **COMPLETE** binary tree.

Question: What are the rules for a **minheap**?

Answer:
1. The value contained by a node is **ALWAYS LESS THAN OR EQUAL TO** the values of the node's children.
2. The tree is a **COMPLETE** binary tree.

The Maxheap

Which of the following are valid **maxheaps**?



Maxheap reminder:

1. The value contained by a node is ALWAYS \geq the values of the node's children.
2. The tree is a COMPLETE binary tree.


Answers:
 1. Upper-left tree: Yes
 2. Lower-left tree: Yes
 3. Middle tree: No - it's not a complete binary tree!
 4. Right tree: Yes

The Maxheap

One thing you'll notice about a maxheap is...

that, by definition, the biggest (highest priority) item

is always at the *root* (top) of the tree!



This means its *easy* to always find the biggest (highest priority) item in just a single step!


Operations on a Maxheap

Allright, now let's see how to **extract an item** from a heap and to **add a new item** to a heap!



Extracting the Biggest Item

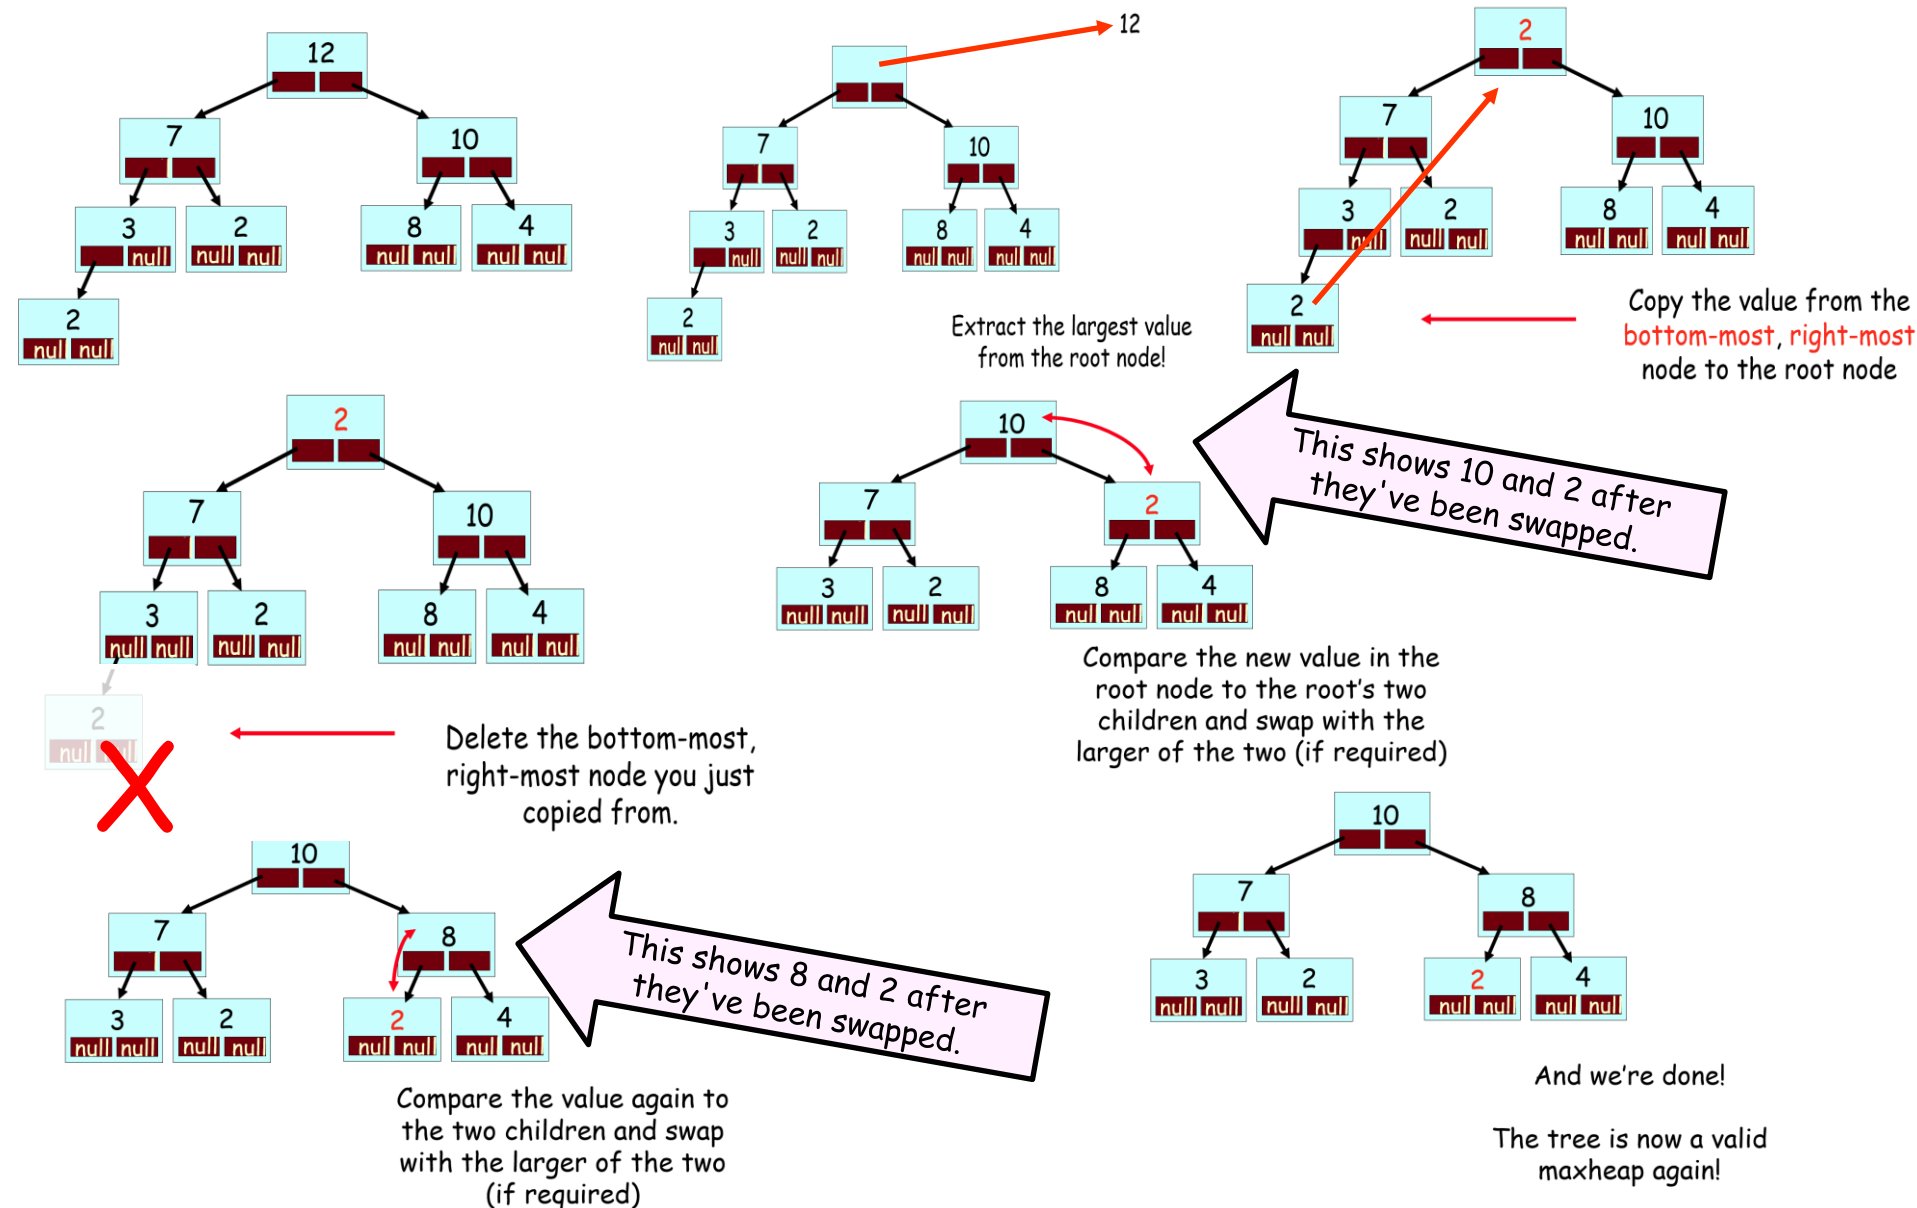
1. If the tree is empty, return error.
2. Otherwise, the top item in the tree is the biggest value. Remember it for later.
3. If the heap has only one node, then delete it and return the saved value.
4. Copy the value from the right-most node in the bottom-most row to the root node.
5. Delete the right-most node in the bottom-most row.
6. Repeatedly swap the just-moved value with the larger of its two children until the value is greater than or equal to both of its children. ("sifting DOWN")
7. Return the saved value to the user.



When we're done, the largest value is on the top again, and the heap is consistent.

Extraction Challenge!


Show the resulting heap after extracting the largest item!



Adding a Node to a Maxheap


(Let's see how to add a value of 9)

1. If the tree is empty, create a new root node & return.
2. Otherwise, insert the new node in the bottom-most, left-most position of the tree (so it's still a complete tree).
3. Compare the new value with its parent's value.
4. If the new value is greater than its parent's value, then swap them.
5. Repeat steps 3-4 until the new value rises to its proper place.



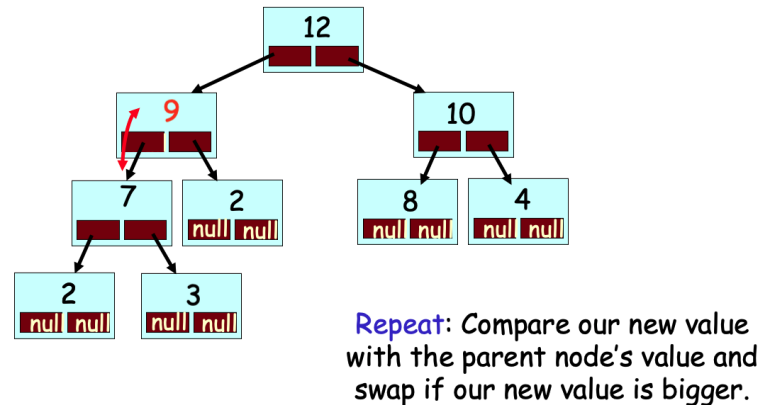
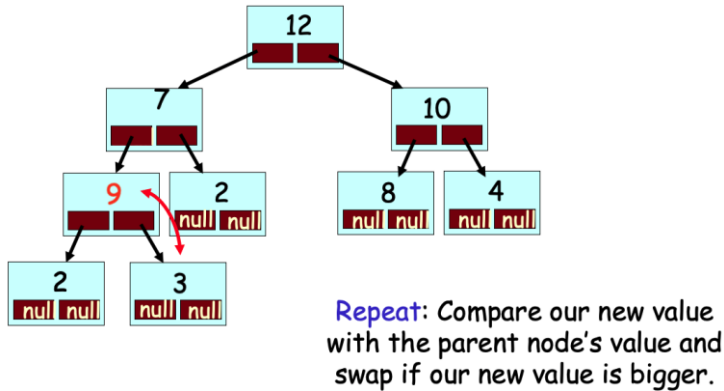


This process is called
"reheapification."

Wot in reheapification?!?!?



Insertion Challenge!

Show the resulting heap after inserting a value of 9!



And now, time for your favorite game!!!!



- > Programming language inventor
- > Serial killer

Implementing A Heap


Question:

What data structure can we use to implement a heap?

How about a classical **binary tree node** with links?

Hmmm... But this has some **challenges**. What are they?

```
struct node
{
    int value;
    node *left, *right;
};
```



1. It's not easy to locate the **bottom-most, right-most** node during **extraction**.
2. It's not easy to locate the **bottom-most, left-most** open spot to insert a new node during **insertion**!
3. It's not easy to locate a **node's parent** to do **reheapification swaps**.


Implementing A Heap

Perhaps there's some better data structure we could use...

Hmmmm. What about an array?

Well, we know that each level of our tree has 2x the number of nodes of the previous level*.

So what if we just copy our nodes a level at a time into an **array**???



* Except for the last level...

Implementing A Heap

So what if we just copy our nodes a level at a time into an array???

Let's see, we can put our root node value in `heap[0]`

And we can put the next two nodes' values into the next two available slots...

And then the next four values in the next four slots...

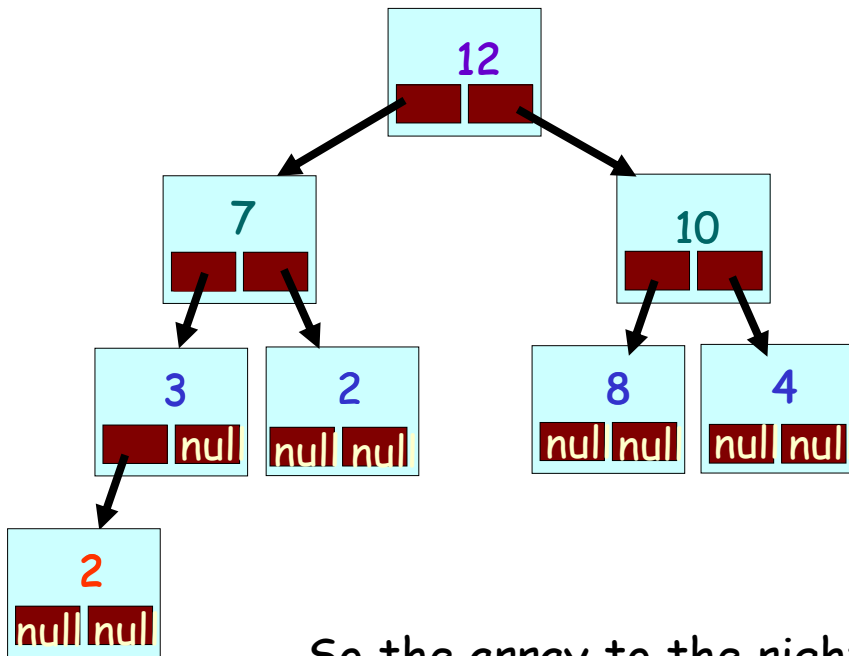
Finally, let's use a simple `int` variable to track how many items are in our heap!

So the array to the right now logically represents the tree on the left! And if we use the array, there's no need to use a node-based tree!


`int count;`
8

`int heap[1000];`

0	12
1	7
2	10
3	3
4	2
5	8
6	4
7	2
8	
9	
10	
11	
12	
13	
...	



Implementing A Heap



int count;
9

int heap[1000];


0	12
1	7
2	10
3	3
4	2
5	8
6	4
7	2
8	1
9	
10	
11	
12	
13	
...	

So what are the properties of our array-based tree?

1. We can always find the root value in `heap[0]`
2. We can always find **bottom-most, right-most node** in `heap[count-1]`
3. We can always find the **bottom-most, left-most empty spot** (to add a new value) in `heap[count]`
4. We can add or remove a node by simply setting `heap[count] = value;` and/or **updating our count!**

Implementing A Heap

int count;
9



Ok, in our array,
how do we locate
the **left** and **right**
children of a node?

Let's consider some examples:

Parent Slot#	Left Child Slot#	Right Child Slot#
0	1	2
2	5	6
3	7	8

Challenge: Come up with a formula to locate a node's children

$$\begin{aligned}\text{leftChild}(\text{parent}) &= 2 * \text{parent} + 1 \\ \text{rightChild}(\text{parent}) &= 2 * \text{parent} + 2\end{aligned}$$


int heap[1000];

0	12
1	7
2	10
3	3
4	2
5	8
6	4
7	2
8	1
9	
10	
11	
12	
13	
14	

Implementing A Heap

```
int count;
    9
```

```
int heap[1000];
```



Let's see, does our formula work?

Consider this node, which is in slot #1 of our array?

$$\text{leftChild}(1) = 2 * 1 + 1 = \text{slot } 3$$

$$\text{rightChild}(1) = 2 * 1 + 2 = \text{slot } 4$$

Our formula appears to work!

Hint: It's of the form

$$\text{leftChild}(\text{parent}) = 2 * \text{parent} + 1$$

$$\text{rightChild}(\text{parent}) = 2 * \text{parent} + 2$$

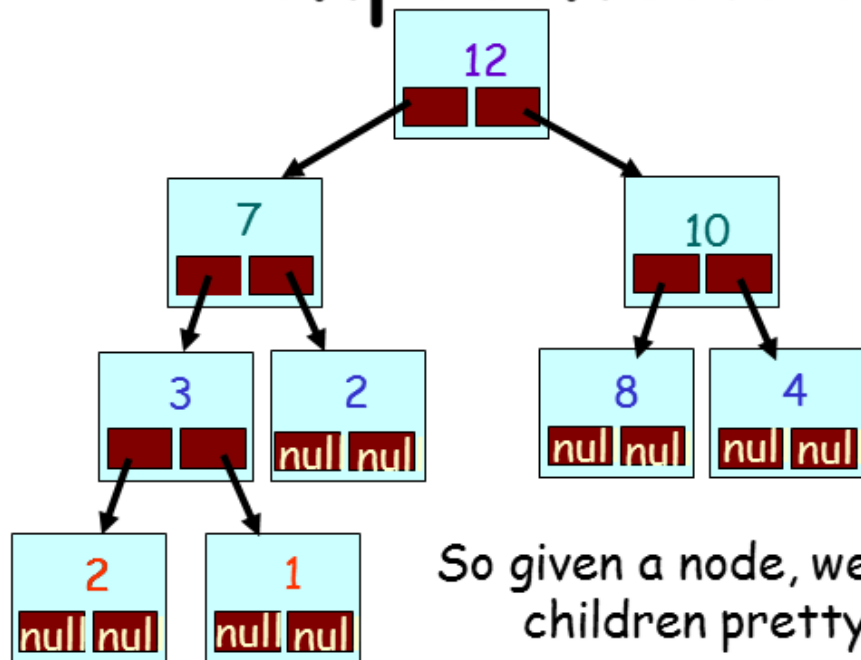
0	12
1	7
2	10
3	3
4	2
5	8
6	4
7	2
8	1
9	
10	
11	
12	
13	

Implementing A Heap

```
int count;
    9
```

```
int heap[1000];
```

0	12
1	7
2	10



So given a node, we can find its two children pretty easily. Cool!

Question: So now how do we find the slot of the parent of some node in our

Answer: Use simple algebra!

$$\frac{\text{child}-1}{2} = \text{parent}$$

And, due to a property of C++ integer division... this formula works equally well for both **left** and **right** children!

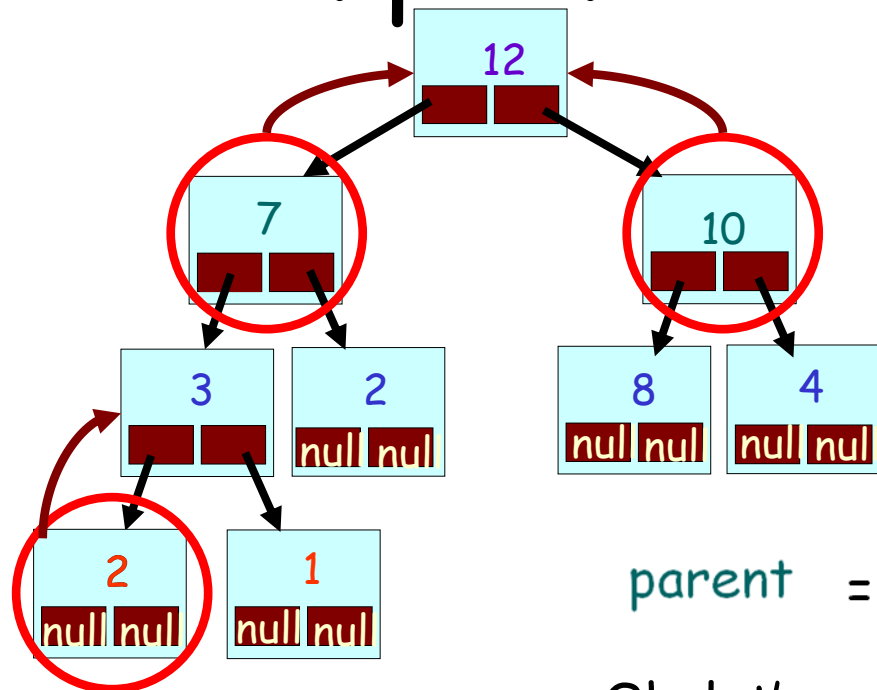
10	
11	
12	
13	
...	

Implementing A Heap

```
int count;
```

9

```
int heap[1000];
```



$$\text{parent} = \frac{\text{child}-1}{2}$$

Ok, let's verify that it works...

The parent of slot #1 is... $(1-1)/2 = 0$

The parent of slot #2 is... $(2-1)/2 = 0$

The parent of slot #7 is... $(7-1)/2 = 3$

0	12
1	7
2	10
3	3
4	2
5	8
6	4
7	2
8	1
9	
10	
11	
12	
13	

Cool stool! So now we know how to **locate the children** of a node, **find the parent** of a node, and **add** and **remove** nodes! ...

Heap in an Array Summary

So, now we know how to store a heap in an array!

Here's a recap of what we just learned :

1. The root of the heap goes in `array[0]`
2. If the data for a node appears in `array[i]`, its children, if they exist, are in these locations:
 Left child: `array[2i+1]`
 Right child: `array[2i+2]`
3. If the data for a non-root node is in `array[i]`, then its parent is always at `array[(i-1)/2]` (Use integer division)

A Heap Helper Class

```
class HeapHelper
{
    HeapHelper()          { num = 0; }
    int GetRootIndex()    { return(0); }
    int LeftChildLoc(int i) { return(2*i+1); }
    int RightChildLoc(int i) { return(2*i+2); }
    int ParentLoc(int i)   { return((i-1)/2); }
    int PrintVal(int i)    { cout << a[i]; }
    void AddNode(int v)    { a[num] = v; ++num;}
private:
    int a[MAX_ITEMS];
    int num;
};
```

Output:

num


5

a[0]	123
[1]	42
[2]	-7
[3]	999
[4]	314

```
main()
{
    HeapHelper a;

    a.AddNode(123);
    a.AddNode(42);
    a.AddNode(-7);
    a.AddNode(999);
    a.AddNode(314);

    int i = GetRootIndex();
    PrintVal(i);
    i = LeftChildLoc(i);
    PrintVal(i);
    i = RightChildLoc(i);
    PrintVal(i);
}
```





Using an Array to Implement a Heap


Ok, so now let's see how to use a simple array to implement a maxheap.

To do so, we'll need to be able to easily do the following operations:

Locate the **root node** of the tree... 

Locate (and delete) the **bottom-most, right-most node** in the tree... 

Add a new node in the **bottom-most, left-most empty position** in the tree... 


Easily locate the **parent** and **children** of any node in the tree... 


Using an Array to Implement a Heap


Ok, so now let's see how to use a simple array to implement a maxheap.

To do so, we'll need to be able to easily do the following operations:

Locate the **root node** of the tree... 

Locate (and delete) the **bottom-most, right-most node** in the tree... 

Add a new node in the **bottom-most, left-most empty position** in the tree... 

Easily locate the **parent** and **children** of any node in the tree... 

Extracting from a Maxheap - The Array Version!

int count;

9

int heap[1000];


1. If the `count == 0` (it's an empty tree), return error.
2. Otherwise, `heap[0]` holds the biggest value. Remember it for later.
3. If the `count == 1` (that was the only node) then `set count=0` and return the saved value.
4. Copy the value from the right-most, bottom-most node to the root node:
`heap[0] = heap[count-1]`
5. Delete the right-most node in the bottom-most row: `count = count - 1`
6. Repeatedly swap the just-moved value with the larger of its two children:
Starting with `i=0`, compare and swap:
`heap[i]` with `heap[2*i+1]` and `heap[2*i+2]`
7. Return the saved value to the user.

0	12
1	7
2	10
3	3
4	2
5	8
6	4
7	2
8	1
9	
10	
11	
12	
13	
...	


Implementing A Heap

Ok, so now let's see how to extract the biggest item from an array-based max-heap!


First, we extract the value from `heap[0]`.
Since this is the largest value in the heap, it should be returned to the user!



Next, we copy the last item in the heap into the root of the tree: `heap[0] = heap[count-1]`




Then we **remove the last node** from the heap by decrementing our **count** variable.




Next we **compare our newly-copied value** to its two children:
`heap[i]` vs. `heap[2*i+1]`
and
`heap[i]` vs. `heap[2*i+2]`

We'll swap it with the larger of the two children, if necessary.



...and we repeat this comparing and swapping process until our copied item is in the right place.



Adding a Node to a Maxheap - The Array Version

1. Insert a new node in the bottom-most, left-most open slot:
 $\text{heap}[\text{count}] = \text{value}$
 $\text{count} = \text{count} + 1;$
2. Compare the new value $\text{heap}[i]$ with its parent's value: $\text{heap}[(i-1)/2]$
3. If the new value is greater than its parent's value, then swap them.
4. Repeat steps 2-3 until the new value rises to its proper place or we reach the top of the array.

int count;

4

int heap[10];


0	10
1	7
2	8
3	3
4	

...


Heap Insertion Challenge

Let's add **12** to our heap.


First, add the new value to the end of our heap array:
`heap[count] = value`
 and then increase our count!




Now, repeatedly compare our new value with it's parent and swap if our new value is bigger:
 if (`heap[i] > heap[(i-1)/2]`)
 then swap



Now, repeatedly compare our new value with it's parent and swap if our new value is bigger:
 if (`heap[i] > heap[(i-1)/2]`)
 then swap



Now, repeatedly compare our new value with it's parent and swap if our new value is bigger:
 if (`heap[1] > heap[(1-1)/2]`)
 then swap



Class Challenge

Show a maxheap and its array after inserting each of the following numbers:

1, 6, 4, 5, 0, 8, 3, 12

1. The root of the binary tree goes in `array[0]`
2. If a node appears in `array[i]`, its children are in these locations:
 Left child: `array[2i+1]`
 Right child: `array[2i+2]`
3. If the data for a non-root node is in `array[i]`, then its parent is always at `array[(i-1)/2]`
(Use integer division)

Class Challenge #2

Now show the maxheap and its array after removing the biggest 2 numbers:

1. The root of the binary tree goes in `array[0]`
2. If a node appears in `array[i]`, its children are in these locations:
 Left child: `array[2i+1]`
 Right child: `array[2i+2]`
3. If the data for a non-root node is in `array[i]`, then its parent is always at `array[(i-1)/2]`
(Use integer division)


Complexity of the Heap

Question: What is the big-oh cost of inserting a new item into a heap?

Every time we insert a new item, we need to keep comparing it with its parent until it reaches the right spot...

Since our tree is a COMPLETE binary tree, if it has N entries, it's guaranteed to be exactly $\log_2(N)$ levels deep.

So in the worst case, we'll have to do $\log_2(N)$ comparisons and swaps of our new value. (This is true whether or not our heap is stored in an array!)



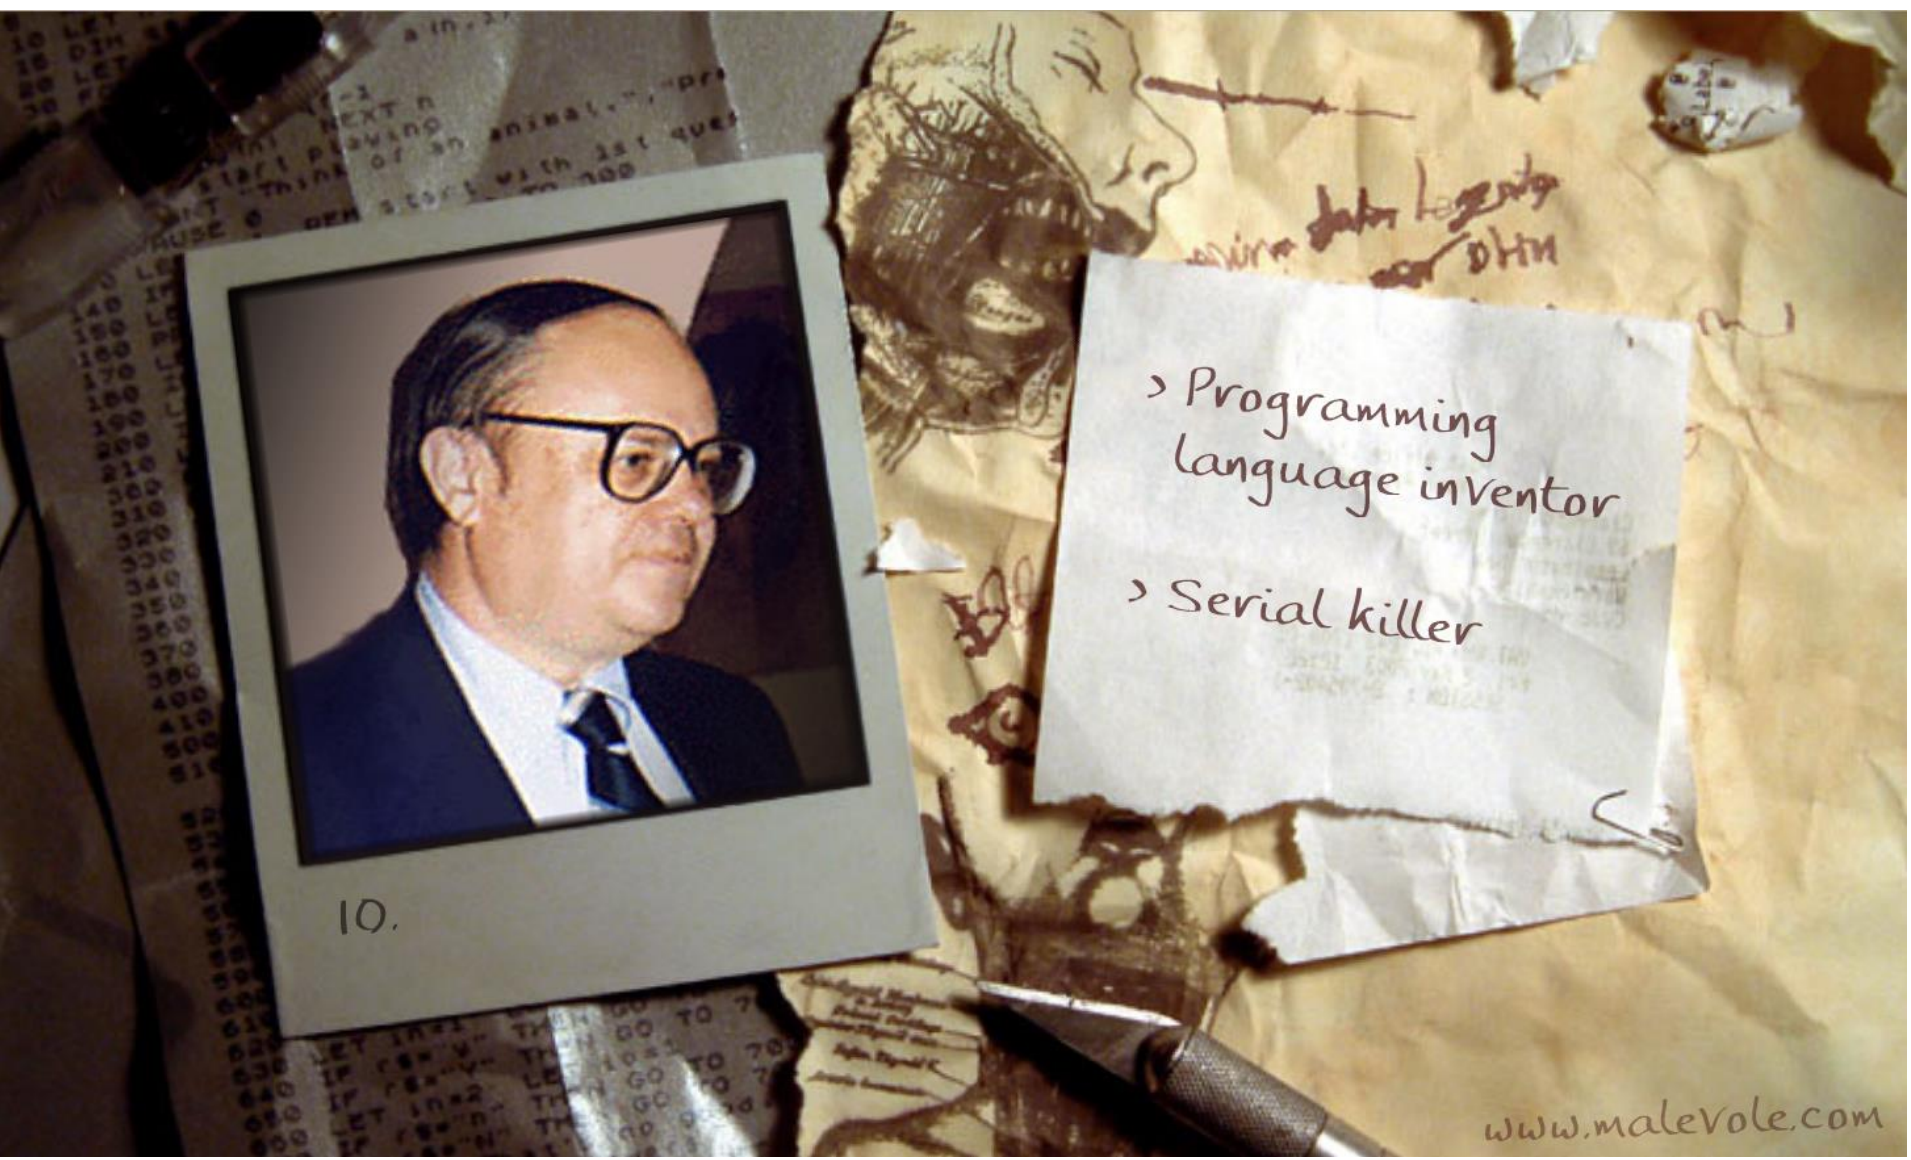
Question: What is the big-oh cost of extracting the maximum/minimum item from a heap?

Just as with heap insertion, when we extract a value we need to bubble an item from the root down the tree.


Since the maximum number of levels in our tree is $\log_2(N)$, the worst case that this requires $\log_2(N)$ swaps.

So inserting and extracting from a heap is $O(\log_2(n))$

Alright, one last time today...



Heapsort



Heapsort

Why should you care?

Errr... Ok, maybe you shouldn't.
Heapsort isn't the most popular
algorithm.

But it'll probably be on your exam, and
I'll feel hurt if you don't listen.

So pay attention!

Why
should
I care?



The Heapsort

Question:


How can we use a **heap** to sort a bunch of items?

Answer:

Here's a "naïve" way to do it...

Given an array of N numbers that we want to sort:

1. **Insert** all N numbers into a new **maxheap**
2. While there are numbers left in the heap:
 - A. **Remove the biggest value** from the heap
 - B. **Place it in the last open slot** of the array



And viola!
Our array is sorted!

7	3	9	1	5
---	---	---	---	---

The Naive Heapsort

Question:

What's the complexity of our simple version of heapsort?

Hints:

What is the cost of **inserting** an item into a maxheap?

What is the cost of **extracting** an item from a maxheap?

How many items must be **inserted** and then **extracted**?

Insertions: N items * $\log_2(N)$ steps per item $\rightarrow N \log_2(N)$ steps

Extractions: N items * $\log_2(N)$ steps per item $\rightarrow N \log_2(N)$ steps

That comes to $2 * N \log_2(N)$ total steps, or $O(N \log_2(N))$.

Not bad! But in fact there's an even **faster way** to use a **heap** to sort an array... Let's see it!

The Efficient Heapsort

In our naïve algorithm, we **took every item** from the array and **inserted** it into a separate, **brand-new maxheap**.

So first we **built a separate maxheap** from scratch, **copying every one of our items over!**

And then we had to **remove each one from our new heap** and **stick them back** into our original array. So SLOW!

Question:

Could we have **avoided creating a whole new maxheap** and moving our numbers back and forth?

Answer:

Yes! That's the way the **"official"** Heapsort works! Let's see!

The Efficient (Official) Heapsort

Let's **update** our original **inefficient algorithm** to turn it into the **efficient (official) version**.

Given an array of N numbers that we want to sort:

1. **Convert our input array into a maxheap**
2. While there are numbers left in the heap:
 - A. Remove the biggest value from the heap
 - B. Place it in the last open slot of the array

Wow! It's that simple? **Yup!**

We're just going to just shuffle the values around in our input array so they become a maxheap.




Then we'll use that maxheap as if it were a separate array like before. Only now everything's in just one array.

Step #1:

Convert your randomly-arranged input array into a **maxheap** by **cleverly shuffling** around the values in the array.

0	3	9	1	5	4	18	21
---	---	---	---	---	---	----	----

21	5	18	3	0	4	9	1
----	---	----	---	---	---	---	---




Let's learn the algorithm to perform this shuffling.

Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree.

By last node, we mean the **bottom-most, right-most node** in the tree. This corresponds to the **last element** in the array.

0	3	9	1	5	4	18	21
---	---	---	---	---	---	----	----




Ok, now here's the algorithm:

for (curNode = lastNode thru rootNode):

Focus on the subtree rooted at curNode.

Think of this subtree as a maxheap.

Keep shifting the top value down until your subtree becomes a valid maxheap.




Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree.

Ok, now here's the algorithm:

0	3	9	1	5	4	18	21
---	---	---	---	---	---	----	----




for (curNode = **lastNode** thru **rootNode**):

Focus on the subtree rooted at curNode.

Think of this subtree as a maxheap.

Keep shifting the top value down until your subtree becomes a valid maxheap.




Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree.

Ok, now here's the algorithm:

0	3	9	1	5	4	18	21
---	---	---	---	---	---	----	----




for (curNode = **lastNode** thru **rootNode**):

Focus on the subtree rooted at curNode.


Think of this subtree as a maxheap.

Keep shifting the top value down until your subtree becomes a valid maxheap.



Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree



Ok, now here's the algorithm:


0	3	9	1	5	4	18	21
---	---	---	---	---	---	----	----

for ($\text{curNode} = \text{lastNode}$ thru rootNode):

Focus on the subtree rooted at curNode .

Think of this subtree as a maxheap.


Keep shifting the top value down until your subtree becomes a valid maxheap.



Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree.

Ok, now here's the algorithm:




for (curNode = **lastNode** thru **rootNode**):

Focus on the subtree rooted at curNode.

Think of this subtree as a maxheap.

Keep shifting the top value down until your subtree becomes a valid maxheap.



Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our current state...

0

for

Now we simply use our normal **heapification** algorithm to turn this subtree into a maxheap.

We'll keep swapping our **root value** down with its larger child until it's bigger than both of its children (or hits a leaf)!

Keep shifting the top value down until your subtree becomes a valid maxheap.

0


(Node):

Node.

Finally - we found a subtree containing more than one node! Let's treat this subtree as if it were a **maxheap**, with a **new value** on top that needs to be sifted down...

curNode


Once we finish, our subtree has been converted into a valid maxheap (see next slide to see 21 and 1 swapped).



Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree.

Ok, now here's the algorithm:




for (curNode = **lastNode** thru **rootNode**):

Focus on the subtree rooted at curNode.

Think of this subtree as a maxheap.

Keep shifting the top value down until your subtree becomes a valid maxheap.

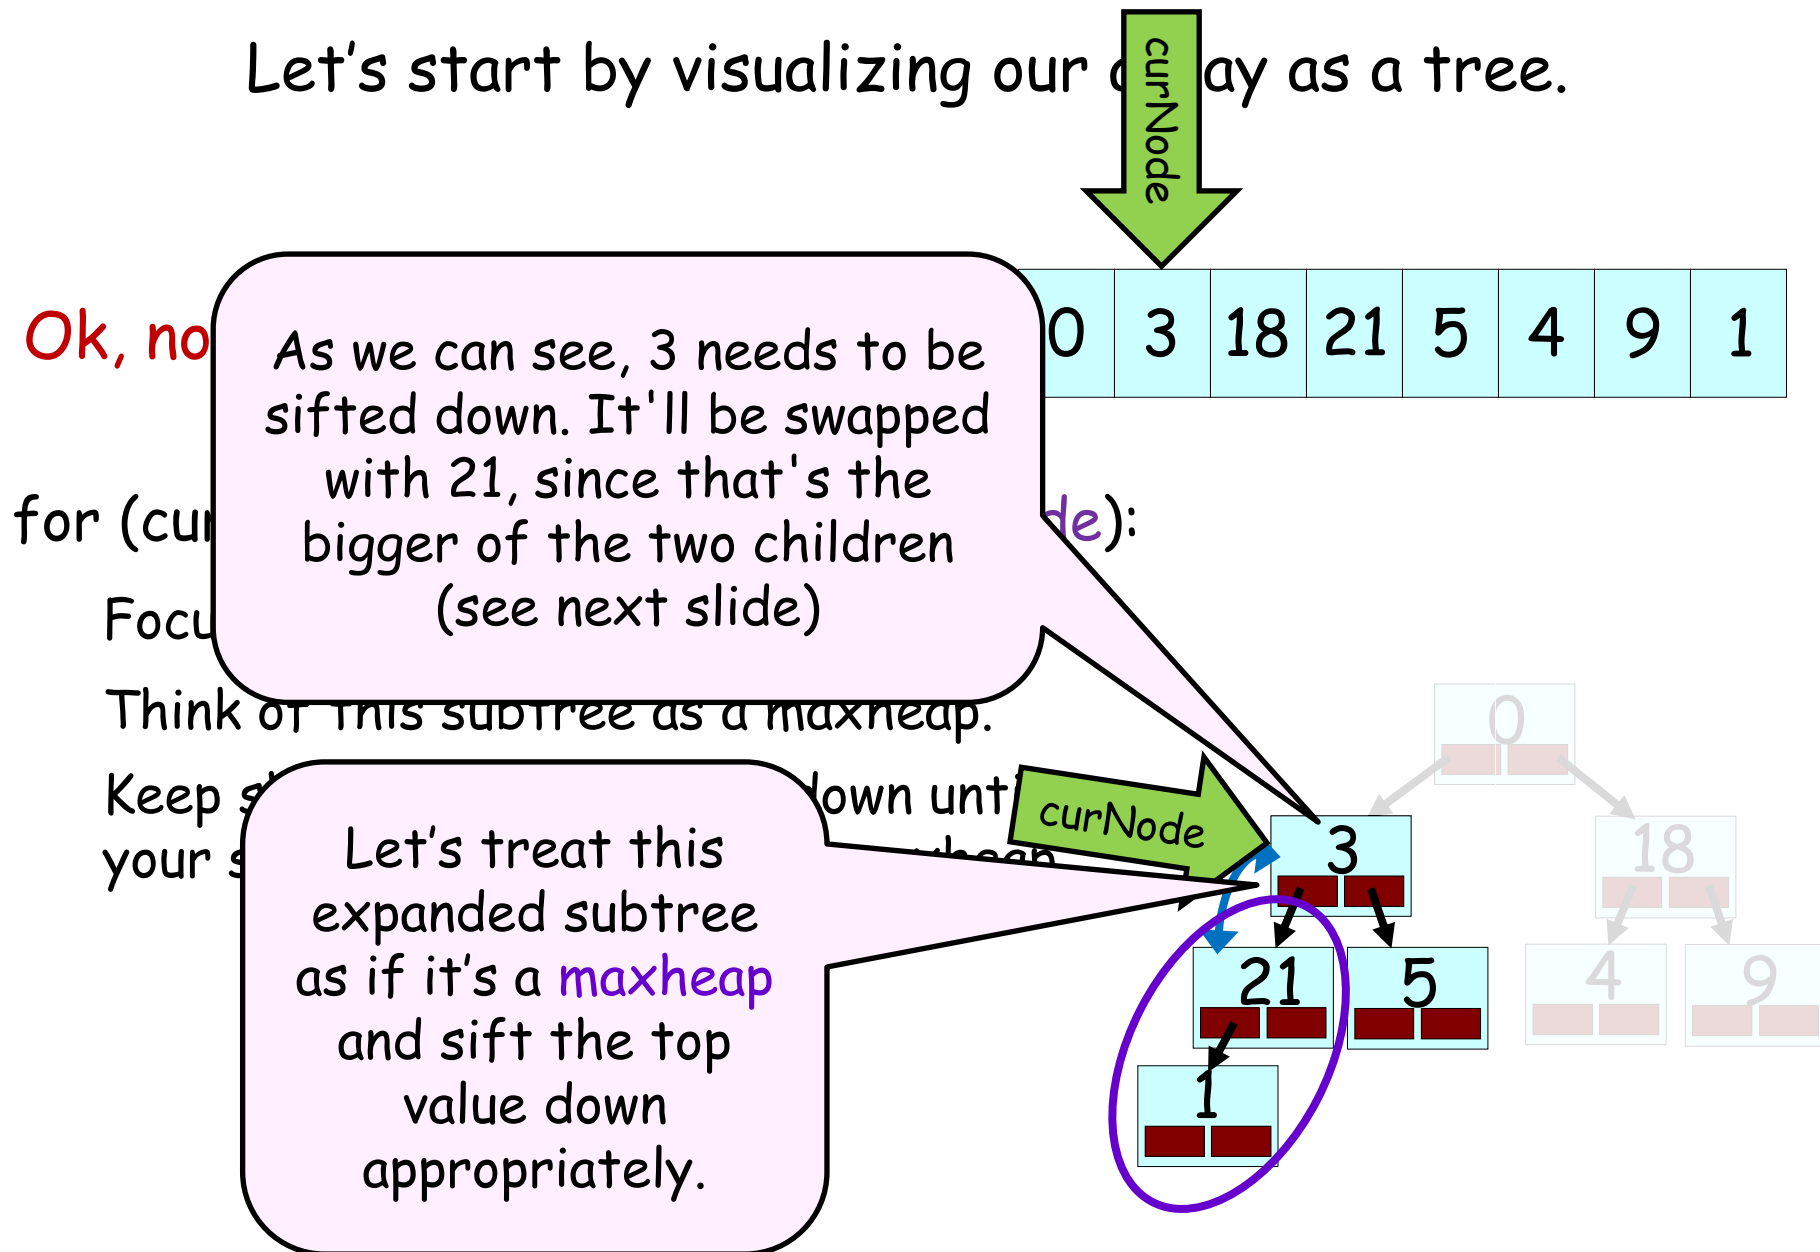
Since 9 is less than 18, we'll bubble it down by swapping with 18 (see next slide).



Again, let's treat this subtree as a **maxheap** with a new value at the top that needs to be sifted down...

Step #1: Convert Your Input Array into a MaxHeap

Let's start by visualizing our array as a tree.



Step #1: Convert Your Input Array into a MaxHeap


Let's start by visualizing our array as a tree.

Essentially what we've done is heapify each sub-tree from the bottom-up.

As we heapify **higher sub-trees**, they rely upon the **lower sub-trees** that were heapified earlier!

Once we've finished heapifying from our **root node**, our entire array will hold a **valid maxheap**!

21	5	18	3	0	4	9	1
----	---	----	---	---	---	---	---



Step #1: Convert Your Input Array into a MaxHeap



There's one more thing to consider!

If you noticed, we wasted a bunch of time looking at single-node sub-trees.

But we only had to reheapify once we reached a sub-tree with at least two nodes.

Wouldn't it be great if we could jump straight to this node to save time?

We can - here's how!



This locates the lowest, right-most node in the tree that has at least one child.

Allowing us to skip all of the single-element trees - that's roughly 50% of all the subtrees!

$\text{startNode} = N/2 - 1$

for (curNode = startNode thru rootNode):



Focus on the subtree rooted at curNode.

Think of this subtree as a maxheap.

Keep shifting the top value down until your subtree becomes a valid maxheap.

This is the complete version of the efficient shuffling algorithm!

Input Array into a MaxHeap



Efficient Heapsort: Step #2


While there are numbers left in the heap:

1. Extract the biggest value from the maxheap and re-heapify (just as we learned about 20 slides ago)


This frees up the last slot in the array
(since the heap now has 1 fewer value in it)

2. Now put the extracted value into this freed-up slot of the array.


A. Array after initial reheapification




B. Array after biggest item extracted from the heap and placed in the last slot of the array.




B. Array after second biggest item extracted from the heap and placed in the 2nd to last slot of the array.




C. Third biggest item processed




D. Fourth biggest item processed



E. Fifth biggest item processed



F. Last item processed



Big-O of Heapsort!

If you think it should be $O(N \log_2 N)$, I did too. ☺ Come to office hours for an explanation.

Step #1:

First we take our N -item array (shown here as a tree) and **convert it into a maxheap**.

We do this from the **bottom up** by converting successively larger **subtrees** into **maxheaps** until the entire tree has been converted.

Step #2:

Then we repeatedly **extract the j^{th} largest item** from the maxheap and place that item **back into the array, j slots from the end**.

Step #1 has a Big-O of $O(N)$.

In other words, we can convert a random array into a maxheap in just $O(N)$ steps!

+

Step #2 has a Big-O of $O(N \log_2 N)$.

Why? Each time we remove an item from the maxheap, it takes $\log_2 N$ steps. We perform this extraction operation N times to sort the entire array.

=

Therefore, Heapsort is $O(N + N \log_2 N)$, which as you know, is just $O(N \log_2 N)$.

HeapSort Challenge

Challenge: Show how to do an in-place heap-sort with the following array of numbers.

5	3	9	6	15	4	11	16
---	---	---	---	----	---	----	----

Step 1: Show the array after each non-leaf is "sifted down" in the array until a valid heap is formed.

Step 2: Show the array after the first 3 items have been removed from the heap and inserted at the end of the array.

(Remember: Sift from $j = N/2 - 1$ down-to $j = 0$)

Challenges

Question: Consider a tri-nary tree, where each node has three children, represented by an array. Given a node i , where can you find its 3 children in the array? Where can you find its parent?

Question: You need to do a fast table search and you also need to print out records in alphabetical order. Which table ADT will you use?

Question: What is the big-oh of traversing a binary search tree?

Answers:

1. $left = 3 * cur + 1, middle = 3 * cur + 2, right = 3 * cur + 3$
2. Use a BST-based map. It's $O(\log N)$ which is the best you can do if you also want the data structure ordered.
3. $O(N)$ - you visit each node once on the way down, once on the way back up from the left, and once on the way back up from the right = $3 * N, O(N)$